

# Introduction to Deep Learning - Homework 1

---

Josh Green

## Problem 1: K-Nearest Neighbors (KNN)

This first problem asks us to use a KNN to recognize handwritten digits.

### Dataset Information

The MNIST dataset can be downloaded using the `download_mnist.py` script. This script generates an `mnist.pkl` file that contains the data converted to numpy arrays. The MNIST dataset consists of:

- `x_train`: A 60,000x784 numpy array containing flattened training images.
- `y_train`: A 1x60,000 numpy array that corresponds to the true label of the corresponding training images.
- `x_test`: A 10,000x784 numpy array where each row contains flattened versions of test images.
- `y_test`: A 1x10,000 numpy array where each component is the true label of the corresponding test images.

There are four compressed files representing the above sets of numpy arrays:

- `train-images-idx3-ubyte.gz`: contains the `x_train` data
- `train-labels-idx1-ubyte.gz`: contains the `y_train` data
- `t10k-images-idx3-ubyte.gz`: contains the `x_test` data
- `t10k-labels-idx1-ubyte.gz`: contains the `y_test` data

### Source Code for `download_mnist.py`

```
import numpy as np
from urllib import request
import gzip
import pickle

filename = [
    ["training_images", "train-images-idx3-ubyte.gz"],
    ["test_images", "t10k-images-idx3-ubyte.gz"],
    ["training_labels", "train-labels-idx1-ubyte.gz"],
    ["test_labels", "t10k-labels-idx1-ubyte.gz"]
]

def download_mnist():
    base_url = "https://oss-ci-datasets.s3.amazonaws.com/mnist/"
    for name in filename:
        print("Downloading " + name[1] + "...")
        request.urlretrieve(base_url + name[1], name[1])
    print("Download complete.")

def save_mnist():
```

```

mnist = {}
for name in filename[:2]:
    with gzip.open(name[1], 'rb') as f:
        mnist[name[0]] = np.frombuffer(f.read(), np.uint8,
offset=16).reshape(-1, 28*28)
for name in filename[-2:]:
    with gzip.open(name[1], 'rb') as f:
        mnist[name[0]] = np.frombuffer(f.read(), np.uint8, offset=8)
with open("mnist.pkl", 'wb') as f:
    pickle.dump(mnist, f)
print("Save complete.")

def init():
    download_mnist()
    save_mnist()

def load():
    with open("mnist.pkl", 'rb') as f:
        mnist = pickle.load(f)
    return mnist["training_images"], mnist["training_labels"],
mnist["test_images"], mnist["test_labels"]

if __name__ == '__main__':
    init()

```

## KNN Information

The `knn.py` script imports the `mnist.pkl` data, looping through a specified data set performing distance calculations. Then, the k-nearest neighbors are calculated to classify the image.

There are two distance calculations that can be used:

- L1 distance is the Manhattan distance, defined as  $D = \sum_{i=1}^n |x_i - y_i|$
- L2 distance is the Euclidian distance, defined as  $D = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$  **NOTE:**  
The `knn.py` script implements the L2 distance formula

## Source Code for `knn.py`

```

import math
import numpy as np
from download_mnist import load
import operator
import time

x_train, y_train, x_test, y_test = load()
x_train = x_train.reshape(60000, 28, 28)
x_test = x_test.reshape(10000, 28, 28)
x_train = x_train.astype(float)
x_test = x_test.astype(float)

def kNNClassify(newInput, dataSet, labels, k):

```

```

    result = []
    for test_sample in newInput:
        distances = []
        for train_image in dataSet:
            distance = np.sqrt(np.sum((test_sample - train_image) ** 2)) #
L2 Distance
            distances.append(distance)

        k_neighbors = np.argsort(distances)[:k]
        k_labels = labels[k_neighbors]
        predicted_label = np.bincount(k_labels).argmax()
        result.append(predicted_label)
    return result

start_time = time.time()
outputlabels = KNNClassify(x_test[0:20], x_train, y_train, 7)
result = y_test[0:20] - outputlabels
result = (1 - np.count_nonzero(result) / len(outputlabels))
print("---classification accuracy for knn on mnist: %s ----" % result)
print("---execution time: %s seconds ----" % (time.time() - start_time))

```

## Output from knn.py

The output from `knn.py` is shown below for a k value of 7 and for 20 images classified.

```

---classification accuracy for knn on mnist: 1.0 ---
---execution time: 4.4463958740234375 seconds ---

```

## KNN Tester Information

The `knn_tester.py` script is a modified version of `knn.py` that runs the knn with both L1 and L2 distances as well as with different k values.

**NOTE:** `knn_tester.py` is classifying 1000 images compared to `knn.py` only classifying 20 images. This was done to get a better representation of the accuracy of the model, and is also why the reported times by the scripts are different.

## knn\_tester.py Script

```

import math
import numpy as np
from download_mnist import load
import operator
import time

x_train, y_train, x_test, y_test = load()
x_train = x_train.reshape(60000,28,28)
x_test = x_test.reshape(10000,28,28)
x_train = x_train.astype(float)

```

```

x_test = x_test.astype(float)

def kNNClassify(newInput, dataSet, labels, k, distance_metric='L2'):
    result=[]
    for test_sample in newInput:
        distances = []

        for train_image in dataSet:
            if distance_metric == 'L2':
                distance = np.sqrt(np.sum((test_sample - train_image) **
2)) #L2 Distance
            else:
                distance = np.sum(np.abs(test_sample - train_image)) #L1
Distance
                distances.append(distance)

            k_neighbors = np.argsort(distances)[:k]
            k_labels = labels[k_neighbors]
            predicted_label = np.bincount(k_labels).argmax()
            result.append(predicted_label)
    return result

# Test both L1 and L2 distances
for distance_type in ['L1', 'L2']:
    print(f"\nTesting {distance_type} Distance:")
    print("=" * 50)

    # Loop through different k values
    for k in range(1, 11):
        start_time = time.time()
        outputlabels = kNNClassify(x_test[0:1000], x_train, y_train, k,
distance_type)
        result = y_test[0:1000] - outputlabels
        accuracy = (1 - np.count_nonzero(result)/len(outputlabels))
        execution_time = time.time() - start_time

        print(f"k={k}:")
        print(f"---classification accuracy for knn on mnist:
{accuracy:.4f} ---")
        print(f"---execution time: {execution_time:.2f} seconds ---")
        print("-" * 50)

```

### Output from knn\_tester.py

```

Testing L1 Distance:
=====
k=1:
---classification accuracy for knn on mnist: 0.9500 ---
---execution time: 123.90 seconds ---
=====

```

```
k=2:
---classification accuracy for knn on mnist: 0.9430 ---
---execution time: 126.50 seconds ---
-----

k=3:
---classification accuracy for knn on mnist: 0.9530 ---
---execution time: 128.91 seconds ---
-----

k=4:
---classification accuracy for knn on mnist: 0.9460 ---
---execution time: 128.22 seconds ---
-----

k=5:
---classification accuracy for knn on mnist: 0.9510 ---
---execution time: 128.85 seconds ---
-----

k=6:
---classification accuracy for knn on mnist: 0.9490 ---
---execution time: 127.98 seconds ---
-----

k=7:
---classification accuracy for knn on mnist: 0.9460 ---
---execution time: 128.89 seconds ---
-----

k=8:
---classification accuracy for knn on mnist: 0.9420 ---
---execution time: 128.22 seconds ---
-----

k=9:
---classification accuracy for knn on mnist: 0.9410 ---
---execution time: 128.45 seconds ---
-----

k=10:
---classification accuracy for knn on mnist: 0.9340 ---
---execution time: 129.16 seconds ---
-----

Testing L2 Distance:
=====
k=1:
---classification accuracy for knn on mnist: 0.9620 ---
---execution time: 149.86 seconds ---
-----

k=2:
---classification accuracy for knn on mnist: 0.9480 ---
---execution time: 150.23 seconds ---
-----

k=3:
---classification accuracy for knn on mnist: 0.9620 ---
---execution time: 150.35 seconds ---
-----

k=4:
---classification accuracy for knn on mnist: 0.9640 ---
---execution time: 149.90 seconds ---
```

```
-----  
k=5:  
---classification accuracy for knn on mnist: 0.9610 ---  
---execution time: 150.04 seconds ---  
-----  
k=6:  
---classification accuracy for knn on mnist: 0.9590 ---  
---execution time: 150.51 seconds ---  
-----  
k=7:  
---classification accuracy for knn on mnist: 0.9620 ---  
---execution time: 152.00 seconds ---  
-----  
k=8:  
---classification accuracy for knn on mnist: 0.9580 ---  
---execution time: 150.28 seconds ---  
-----  
k=9:  
---classification accuracy for knn on mnist: 0.9520 ---  
---execution time: 149.62 seconds ---  
-----  
k=10:  
---classification accuracy for knn on mnist: 0.9540 ---  
---execution time: 150.26 seconds ---  
-----
```

## Observations on L1 and L2 Distances and Accuracy as K Changes

After running the `knn_tester.py` script, it can be observed that the accuracy improved and the runtime increased when using the L2 distance formula. From my small sample size, the reported times don't show a correlation between accuracy and an increased K value. As K increases, we initially see the accuracy increase, but we eventually get to a point where we start considering neighbors that are unrelated to an image we are looking at, which can account for the lower accuracy as K gets closer to 10.

---

## Problem 2: Linear Classifier

The second problem asks us to train a linear classifier to recognize handwritten digits.

The `linear_classifier.py` script implements a linear classifier to accomplish handwriting recognition. This linear classifier uses "Cross Entropy" for the loss function and employs "Random Search" to find the parameters  $W$ . Afterwards, the accuracy of the linear classifier is tested using the MNIST testing set.

Source Code for `linear_classifier.py`

```
import numpy as np  
from download_mnist import load  
from sklearn.preprocessing import OneHotEncoder  
import time
```

```
# Load MNIST dataset
x_train, y_train, x_test, y_test = load()
x_train = x_train.reshape(60000, -1) / 255.0
x_test = x_test.reshape(10000, -1) / 255.0

# One-hot encode labels
encoder = OneHotEncoder(sparse_output=False)
y_train_onehot = encoder.fit_transform(y_train.reshape(-1, 1))
y_test_onehot = encoder.transform(y_test.reshape(-1, 1))

# Hyperparameters
num_classes = 10
num_features = x_train.shape[1]
learning_rate = 0.1
num_epochs = 100
batch_size = 128

# Initialize weights
W = np.random.randn(num_features, num_classes) / np.sqrt(num_features)
start_time = time.time()

# Training loop
for epoch in range(num_epochs):
    # Shuffle training data
    indices = np.random.permutation(len(x_train))
    x_train_shuffled = x_train[indices]
    y_train_shuffled = y_train_onehot[indices]

    # Mini-batch training
    for i in range(0, len(x_train), batch_size):
        batch_x = x_train_shuffled[i:i + batch_size]
        batch_y = y_train_shuffled[i:i + batch_size]

        # Forward pass
        scores = batch_x @ W
        exp_scores = np.exp(scores - np.max(scores, axis=1,
keepdims=True))
        probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)

        # Backward pass
        dscores = probs - batch_y
        dW = batch_x.T @ dscores

        # Update weights
        W -= learning_rate * dW / batch_size

    # Evaluate on test set
    if (epoch + 1) % 10 == 0:
        test_scores = x_test @ W
        test_preds = np.argmax(test_scores, axis=1)
        accuracy = np.mean(test_preds == y_test)
        print(f'Epoch {epoch + 1}, Test accuracy: {accuracy:.4f}')

print("---execution time: %s seconds ---" % (time.time() - start_time))
```

Results for `linear_classifier.py`

```
Epoch 10, Test accuracy: 0.9195
Epoch 20, Test accuracy: 0.9228
Epoch 30, Test accuracy: 0.9230
Epoch 40, Test accuracy: 0.9255
Epoch 50, Test accuracy: 0.9248
Epoch 60, Test accuracy: 0.9252
Epoch 70, Test accuracy: 0.9243
Epoch 80, Test accuracy: 0.9255
Epoch 90, Test accuracy: 0.9251
Epoch 100, Test accuracy: 0.9259
---execution time: 13.739806652069092 seconds ---
```